



# Terraform



**Sabreen Salama**

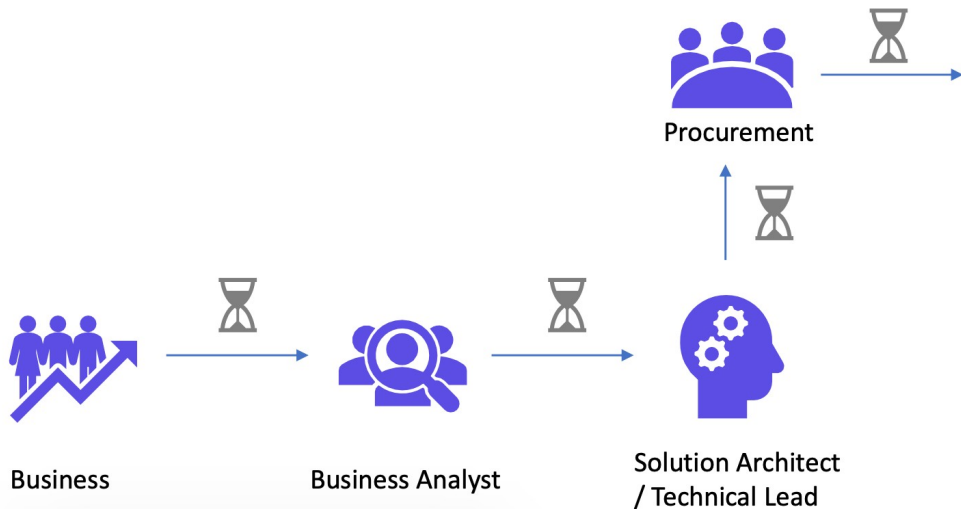
[Linkedin](#)



## Content

---

- Intro to IAC
- IaC Tools
- What is terraform?
- Terraform basics
- Terraform state
- Terraform remote state
- Terraform Provisioner
- Working with terraform
- Terraform import
- Terraform modules
- CloudFormation



Slow Deployment



Expensive



Limited Automation



Human Error



Inconsistency

Infrastructure Team



Field Engineers



System/ NW  
Administrators




Storage Admins



Backup Admins



Application team





## Before IAC?

---

- In the past, managing IT infrastructure was a hard job. System administrators had to manually manage and configure all of the hardware and software that was needed for the applications to run and this takes more time and efforts some human errors due to inconsistency



## Infrastructure as code

---

Allow us to create and manage infrastructure with configuration files rather than through a graphical user interface or manual handling scripts. These files can be versioned, reused, and shared.

# Types of IAC Tools

## Configuration Management



## Server Templating



## Provisioning Tools





## Why using infrastructure as code?

---

- ⦿ **Consistence:** you guarantee the same configurations will be deployed over and over, without discrepancies
- ⦿ **Lower cost**
- ⦿ **Speed:** enables you to quickly set up your complete infrastructure by running a script
- ⦿ **Accountability:** you can version IaC configuration files like any source code file configuration



## Terraform

- ◎ Open source , Free , created by HashiCorp's infrastructure as code tool . It lets you define resources and infrastructure in human-readable, declarative configuration files, and manages your infrastructure's lifecycle to many providers, can be applied on private , public cloud , and on prem.
  
- ◎ Why using terraform?
  - Terraform can manage infrastructure on multiple cloud platforms
  - The human-readable configuration language helps you write infrastructure code quickly
  - Terraform's state allows you to track resource changes throughout your deployments.
  - You can commit your configurations to version control to



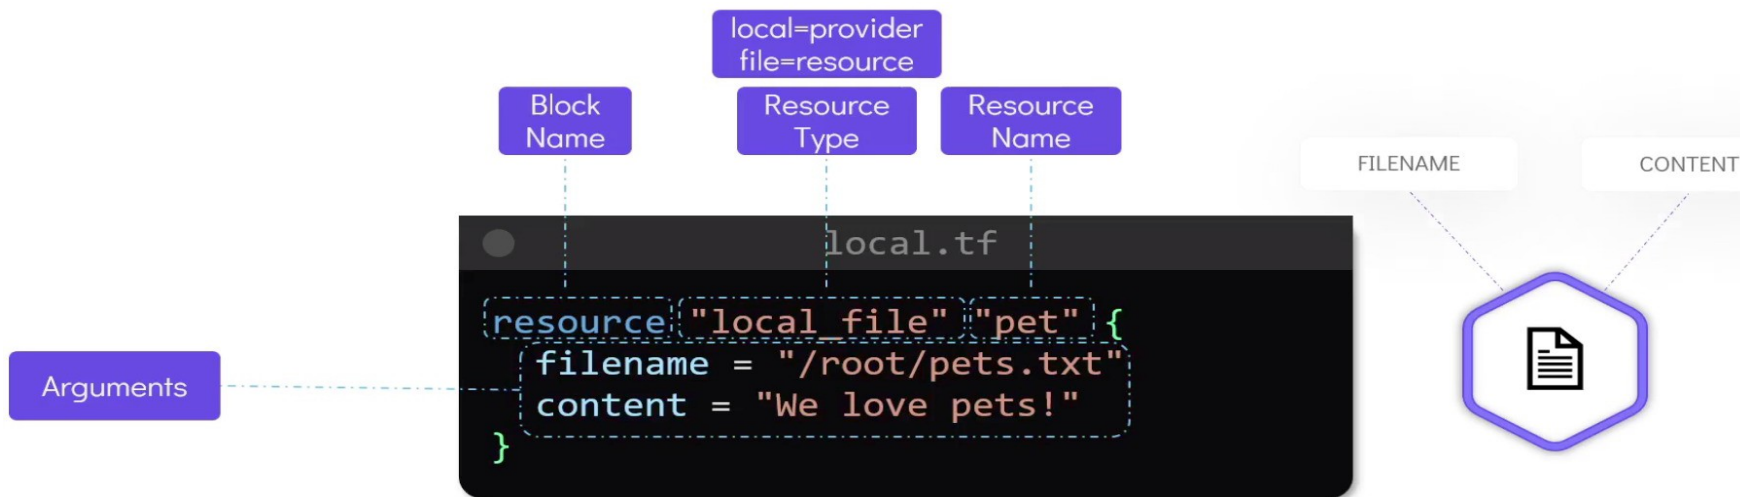


1

# Terraform insatll




# HCL





## Terraform Lifecycle





## provider

☉ A provider is a plugin that interacts with the various APIs required to create, update, and delete various resource



```
terraform {  
  required_providers {  
    aws = {  
      source = "hashicorp/aws"  
      version = "~> 4.0"  
    }  
  }  
}  
  
# Configure the AWS Provider  
provider "aws" {  
  region = "us-east-1"  
}
```



## Contd: Multiple provider

```
resource "local_file" "foo" {  
    content = "Hello from ITI"  
    filename = "./iti-terraform"  
}  
  
resource "aws_s3_bucket" "test-s3" {  
    bucket = "my-tf-test-bucket"  
  
    tags = {  
        Name          = "My bucket"  
        Environment = "Dev"  
    }  
}
```



## Input variables

- it is a key-value pair used as parameters to input values at run time to enable reusability to our code.

```
variables.tf > variable "content" > de
1  variable "content" {
2      default= "ITI"
3  }
```

```
local.tf > resource "local_file" "foo" > content
1  resource "local_file" "foo" {
2      content = var.content
3      filename = "./iti-terraform"
4  }
```

## Variable Definition Precedence

1

>\_

```
$ terraform apply -var "filename=/root/best-pet.txt"
```

2

variable.auto.tfvars

```
filename = "/root/mypet.txt"
```

3

terraform.tfvars

```
filename = "/root/pets.txt"
```

4

>\_

```
$ export TF_VAR_filename="/root/cats.txt"
```



## Dependency of resources

---

- Resource attribute or reference attribute: to use output of a resource as an input for another resource





## Reference Attributes

```
resource "local_file" "foo" {  
    content  = var.content  
    filename = "./iti-terraform"  
}  
  
resource "aws_s3_bucket" "test-s3"  
    bucket = local_file.foo.filename  
  
    tags = {  
        Name           = "My bucket"  
        Environment    = "Dev"  
    }  
}
```



## Contd Dependency of resources

- Explicit dependency:  
Make use of all configuration files without making use of reference

```
.tf > resource "aws_s3_bucket" "test-s3" > tag
resource "local_file" "foo" {
    content = var.content
    filename = "./iti-terraform"
    depends_on = [
        aws_s3_bucket.test-s3
    ]
}
resource "aws_s3_bucket" "test-s3" {
    bucket = "my-tf-test-bucket"

    tags = {
        Name = "My bucket"
        Environment = "Dev"
    }
}
```



## Outputs

- Return attributes for a resource to print it to the screen or to feed it to configuration tool as ansible

```
output "instance_ip_addr" {  
    value = aws_instance.server.private_ip  
}
```

Copy 



## Terraform state

---

It is a json data structure file that maps a real world infra to the resource definition in configuration files

Purpose of state file:

- Mapping resources into real world resources
- Tracking metadata such as dependency
- Collaboration



## Terraform remote state

```
terraform {  
  backend "s3" {  
    # Replace this with your bucket name!  
    bucket      = "terraform-up-and-running-state"  
    key         = "stage/data-stores/mysql/terraform.tfstate"  
    region      = "us-east-2"  
  
    # Replace this with your DynamoDB table name!  
    dynamodb_table = "terraform-up-and-running-locks"  
    encrypt         = true  
  }  
}
```



## Terraform commands

---

- Terraform validate
- Terraform fmt
- Terraform show
- Terraform providers
- Terraform output
- Terraform refresh
- Terraform graph



## Terraform commands

---

- Terraform state list
- Terraform state mv
- Terraform state rm
- Terraform state show



## Lifecycle rules

The default when we run terraform apply the old resource deleted and then the new one will be created what about not deleting the old ones or creating first then deleting

```
resource "kubernetes_namespace" "test" {  
  lifecycle {  
    create_before_destroy = true  
  }  
  metadata {  
    name = var.ns_name  
  }  
}
```





## Datasources

Allow Terraform to read attributes from a resources that were created manually outside of terraform

```
data "aws_ami" "example" {  
    most_recent = true  
  
    owners = ["self"]  
    tags = {  
        Name    = "app-server"  
        Tested  = "true"  
    }  
}
```



## Resource VS DataSource

```
resource "aws_instance" "web" {  
    ami          = data.aws_ami.web.id  
    instance_type = "t1.micro"  
}
```



## Meta argument

To change the behaviour of a resource as depends\_on and lifecycle

Now we will deal with count and for\_each

```
resource "aws_iam_user" "the-accounts" {  
  for_each = toset( ["Todd", "James", "Alice", "Dottie"] )  
  name     = each.key  
}
```

Copy 



## Count VS for\_each

```
resource "aws_internet_gateway" "example" {  
  # One Internet Gateway per VPC  
  
  for_each = aws_vpc.example  
  
  # each.value here is a full aws_vpc object  
  vpc_id = each.value.id  
}
```

```
resource "aws_instance" "server" {  
  count = 4 # create four similar EC2 instances  
  
  ami           = "ami-a1b2c3d4"  
  instance_type = "t2.micro"  
  
  tags = {  
    Name = "Server ${count.index}"  
  }  
}
```



## Terraform import

```
main.tf

resource "aws_instance" "webserver-2" {
  # (resource arguments)
}
```

```
>_

$ terraform import aws_instance.webserver-2 i-026e13be10d5326f7

aws_instance.webserver-2: Importing from ID "i-026e13be10d5326f7"...
aws_instance.webserver-2: Import prepared!
  Prepared aws_instance for import
aws_instance.webserver-2: Refreshing state... [id=i-026e13be10d5326f7]

Import successful!

The resources that were imported are shown above. These resources are now in
your Terraform state and will henceforth be managed by Terraform.
```



## remote-exec Provisioner

The remote-exec provisioner invokes a script on a remote resource after it is created. This can be used to run a configuration management tool, bootstrap into a cluster, etc

```
resource "aws_instance" "web" {  
  connection {  
    type      = "ssh"  
    user      = "root"  
    password  = var.root_password  
    host      = self.public_ip  
  }  
  
  provisioner "remote-exec" {  
    inline = [  
      "chmod +x /tmp/script.sh",  
    ]  
  }  
}
```



## local-exec Provisioner

The **local-exec** provisioner invokes a local executable after a resource is created. This invokes a process on the machine running Terraform, not on the resource.

```
resource "aws_instance" "web" {  
  # ...  
  
  provisioner "local-exec" {  
    command = "echo ${self.private_ip} >> private_ips.txt"  
  }  
}
```



## Terraform modules

---

A module is a container for multiple resources that are used together. Modules can be used to create lightweight abstractions, so that you can describe your infrastructure in terms of its architecture, rather than directly in terms of physical objects.



child module "server"

```
resource "type" "vm" {  
  network_id = var.network_id  
  . . .  
  . . .  
}
```

```
variable "network_id" {  
}
```

child module "network"

```
resource "type" "network" {  
  argument = value  
  argument = value  
  argument = value  
  . . .  
}
```

```
output "network_id" {  
  value = type.network.id  
}
```

 I have a solution!

root module



## Terraform Workspace

---

Each Terraform configuration has an associated [backend](#) that defines how Terraform executes operations and where Terraform stores persistent data, like [state](#).

The persistent data stored in the backend belongs to a workspace. The backend initially has only one workspace containing one Terraform state associated with that configuration.



# Terraform Workspace

To list all spaces:

**terraform workspace list**

For creating new space:

**terraform workspace new  
<space-name>**


```
resource "aws_instance" "example" {  
  count = "${terraform.workspace == "default" ? 5 : 1}"  
  
  # ... other arguments  
}
```

Copy

Another popular use case is using the workspace name as part of naming or tagging behavior:

```
resource "aws_instance" "example" {  
  tags = {  
    Name = "web - ${terraform.workspace}"  
  }  
  
  # ... other arguments  
}
```

Copy



# Thanks!

*Any questions ?*